YOULearn.com: A Case Study in Cloud-Based Data Management

By: Rey Anthony Rodriguez | 000931509

DATA040 – 008

Table of Contents

YOULearn.com
Your Online University

I. Introduction

Online education platforms manage large volumes of data, from user profiles to course content and platform usage. YOULearn.com processes structured and unstructured data for learners, instructors, and administrators. To handle this efficiently, a scalable, cloud-native data architecture was designed using Microsoft Azure services. This case study documents the solution, which supports real-time insights, personalized learning, and operational visibility.

II. Objectives

- Build a scalable cloud system to integrate diverse data sources.
- Provide real-time insights for students, instructors, and administrators.
- Ensure system reliability, performance, and security through end-to-end monitoring.

III. Architecture Overview

The system collects data from different sources, processes it, and delivers key outputs.

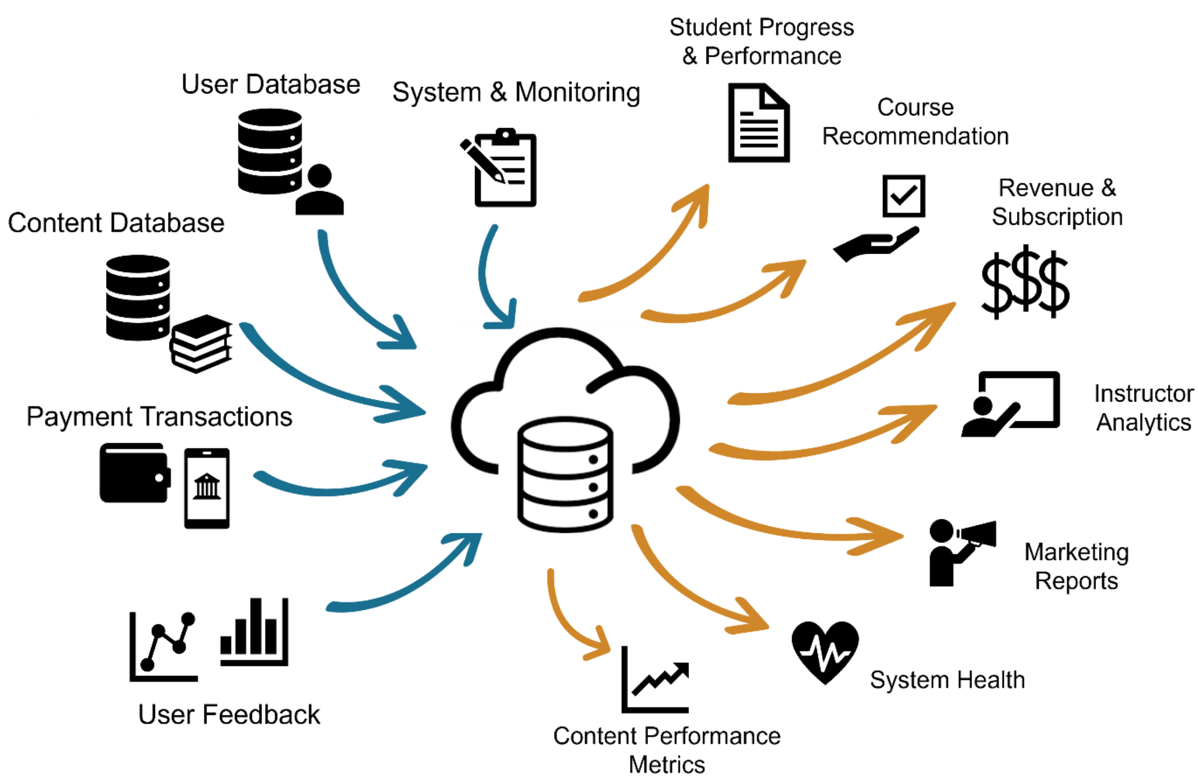

IV. Data Sources

Source	Type	Location	Description
User Database	Structured, SQL tables.	SQL DB	Stores student and instructor profiles, enrollments, and progress.
Course Content Storage	Unstructured, multimedia files.	Microsoft OneDrive	Stores videos, documents, quizzes, and interactive content.
Payment Transactions	Structured, SQL tables.	SQL DB	Stores subscription payments, refunds, and invoices.
User Engagement Data	Structured, SQL tables.	SQL DB	Tracks course completion, time spent, quiz scores, and activity.
System Logs	Semi-structured, log files.	AWS CloudWatch, ELK Stack	Captures system events, errors, and performance metrics.

V. Data Outputs

Output	Delivery	Used By	Example
Student Reports & Certifications	Auto-generated reports and dashboards.	Students, instructors, administrators	A course completion certificate.
Course Recommendations	AI-powered recommendation engine.	Students.	"You might like 'Python for Beginners.'"
Instructor Dashboards	Web-based analytics dashboard.	Instructors.	Course completion rates and student engagement.
Marketing Insights	AI-driven reports.	Marketing teams.	User retention trends.
Revenue Reports	Financial dashboards and reports.	Business and finance teams.	Monthly subscription revenue breakdown.
Content Performance Metrics	Reports and analytics dashboards.	Instructors, content teams.	Dropout rates for each lecture.
System Health Reports	Real-time monitoring dashboards.	IT and DevOps teams.	Server uptime and error logs.

VI. Cloud Architecture

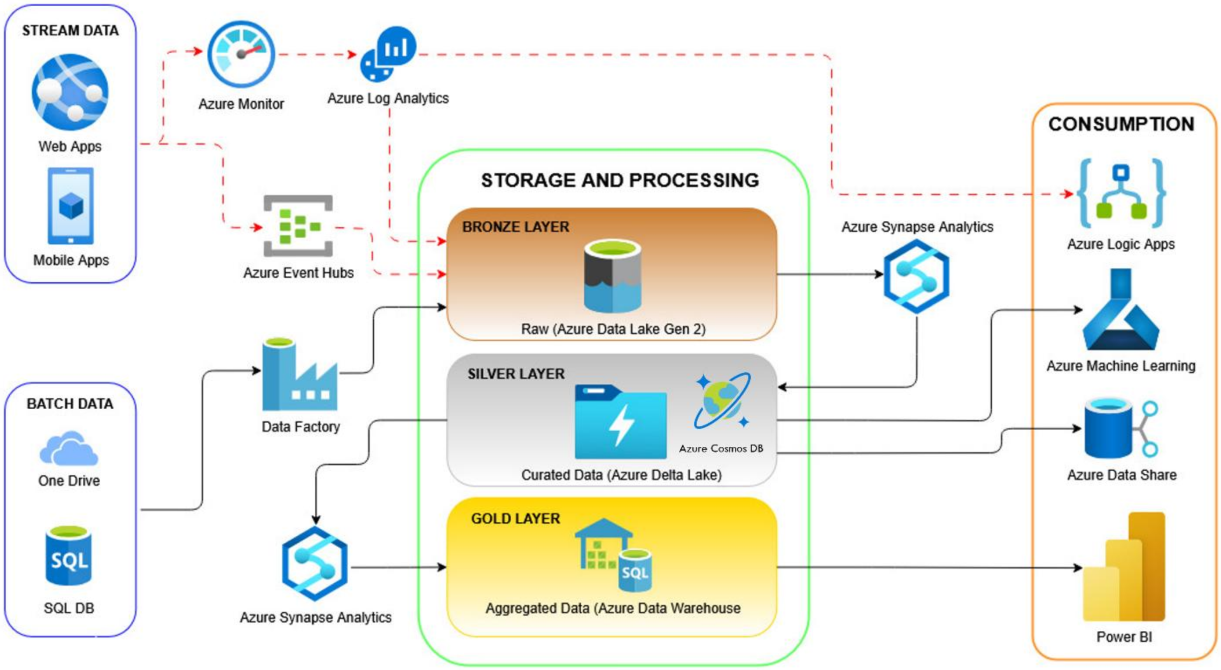

The data architecture is built around a Lakehouse model using the Bronze-Silver-Gold framework. It supports both batch and real-time data sources. The system consolidates raw, curated, and aggregated data for different use cases—from real-time dashboards to machine learning.

I added Azure Monitor + Log Analytics to track system health and logs in real-time. Azure Data Share enables secure internal or partner-level data collaboration. I've tightly integrated Azure ML for AI-powered personalization—like smart course recommendations.

Finally, Power BI delivers live, interactive dashboards for leadership and faculty decision-making.

VII. Pipeline Design

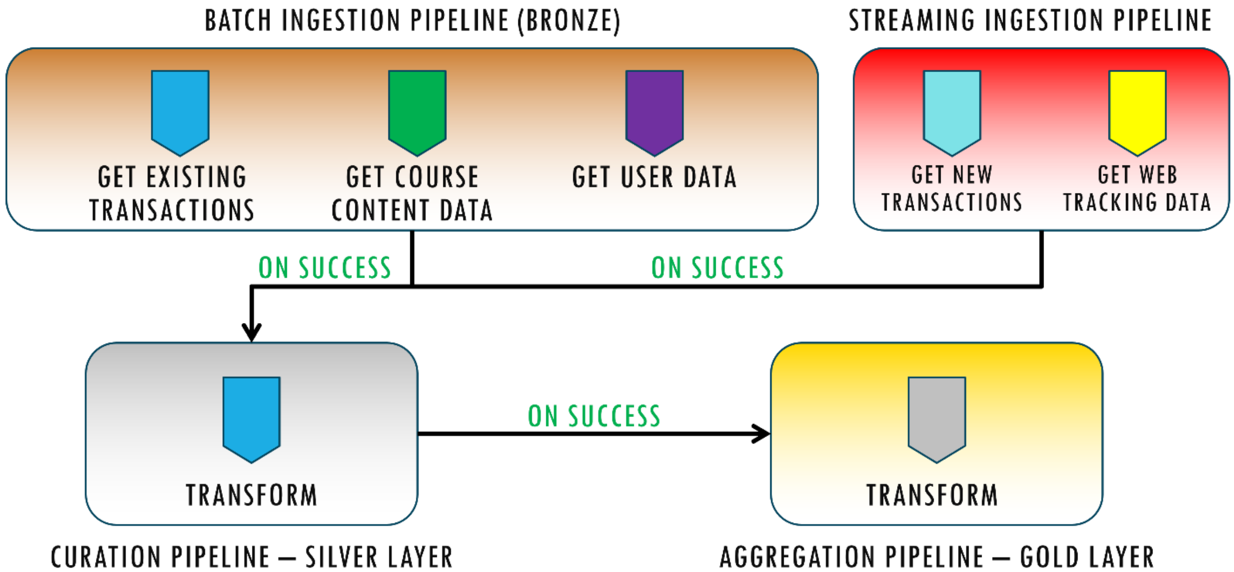

VII.I Batch Ingestion:

I use **Azure Data Factory** to perform a one-time migration of existing data from the following sources:

- User Database
- User Feedback
- Content Database (from OneDrive)
- SQL Database (historical transactions)

These datasets are extracted, loaded into staging, and then dumped into the Bronze Layer. The goal is to preserve the raw form of data for audit, reprocessing, or recovery needs.

VII.II Streaming Ingestion:

For real-time operations and insights, we capture new and dynamic data from various user and system activities such as:

- User registrations
- Payment transactions
- User interactions (clicks, course views)
- Instructor video uploads

VII.III Silver Layer: Curation & Validation

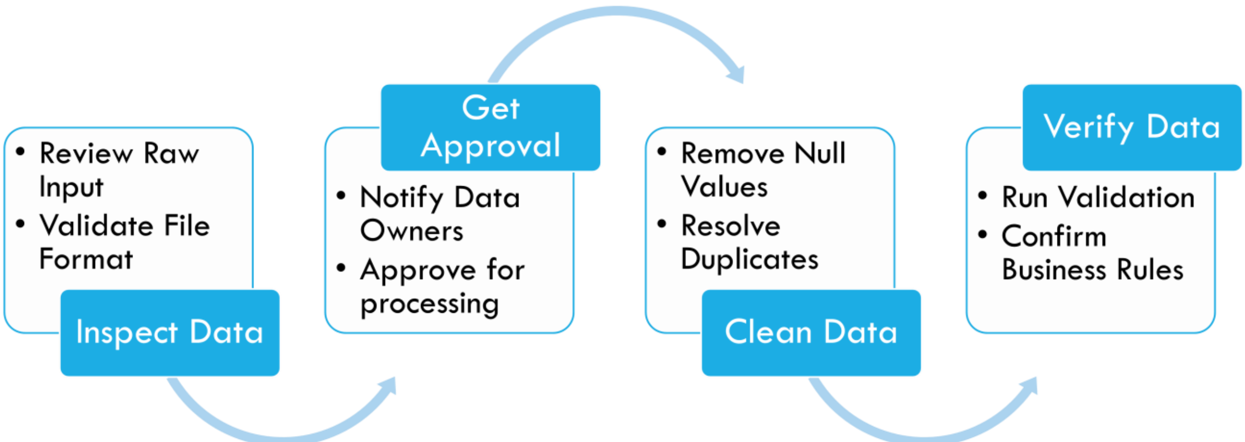

1. **Inspect Data:**
    1. Datasets are profiled to detect anomalies, missing values, duplicate records, and inconsistencies.
    2. For instance, verifying if every payment has a matching user and course, or ensuring video uploads have metadata.
2. **Get Approval:**
    1. Before mass transformation, curated data goes through a **validation checkpoint**.
    2. This is where domain experts (admin, data owners) confirm whether the data is trustworthy for business use.
3. **Clean Data:**
    1. This involves removing nulls, correcting types, applying business logic (e.g., converting durations, mapping user activity), and formatting timestamps.
4. **Verify Data:**
    1. A final verification step is done using SQL queries and data quality checks.
    2. Example: Ensuring every user_id is unique and active, or that feedback is linked to an existing course.

The Silver Layer forms the **foundation for analytics**, enabling business teams to access clean, structured data across different domains: users, content, finance, and performance.

VII.IV Gold Layer : Aggregation & Business KPIs

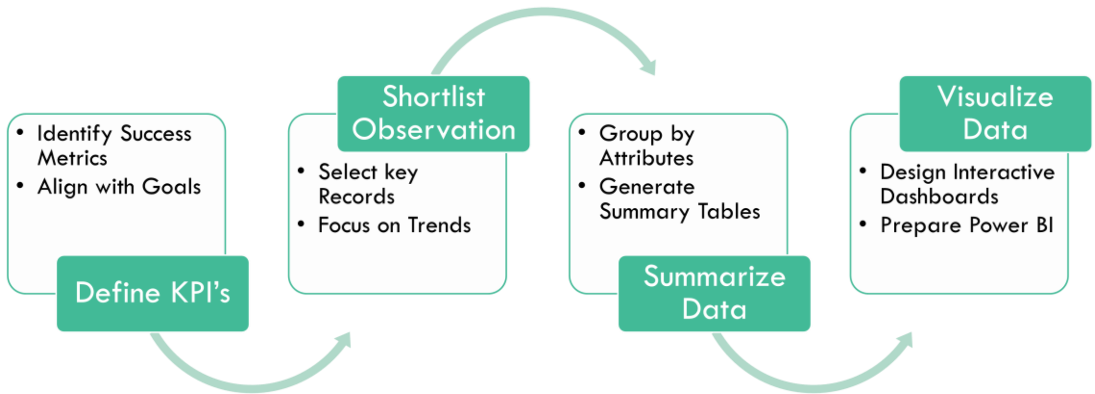

The Gold Layer transforms curated datasets into **insightful, business-ready views**, typically used by executive dashboards and reporting tools.

Aggregation Workflow:

1. **Define KPIs:** Collaborate with stakeholders to identify critical metrics such as:

   Student Progress Rate | Instructor Upload Frequency | Course Completion Rate | Monthly Revenue & Subscriptions

2. **Shortlist Observations:** The clean data from the Silver Layer is filtered to extract only relevant observations per KPI. For example, filtering user activity logs to only include learning sessions longer than 5 minutes.

3. **Summarize Data:** Aggregations are performed using **Azure Synapse Dedicated SQL Pool**, transforming millions of records into compact summaries:
   - AVG, SUM, COUNT, MIN/MAX
   - Grouped by time, course ID, region, etc.

4. **Visualize Data:** Aggregated datasets are connected to **Power BI**, enabling dashboards and drilldowns:

   Content Performance Matrix | Revenue Reports | Instructor Analytics | Marketing Funnel Visuals

The Gold Layer ensures that decision-makers always have access to **reliable, high-quality data views**.

VII.V PIPELINE APPROACH

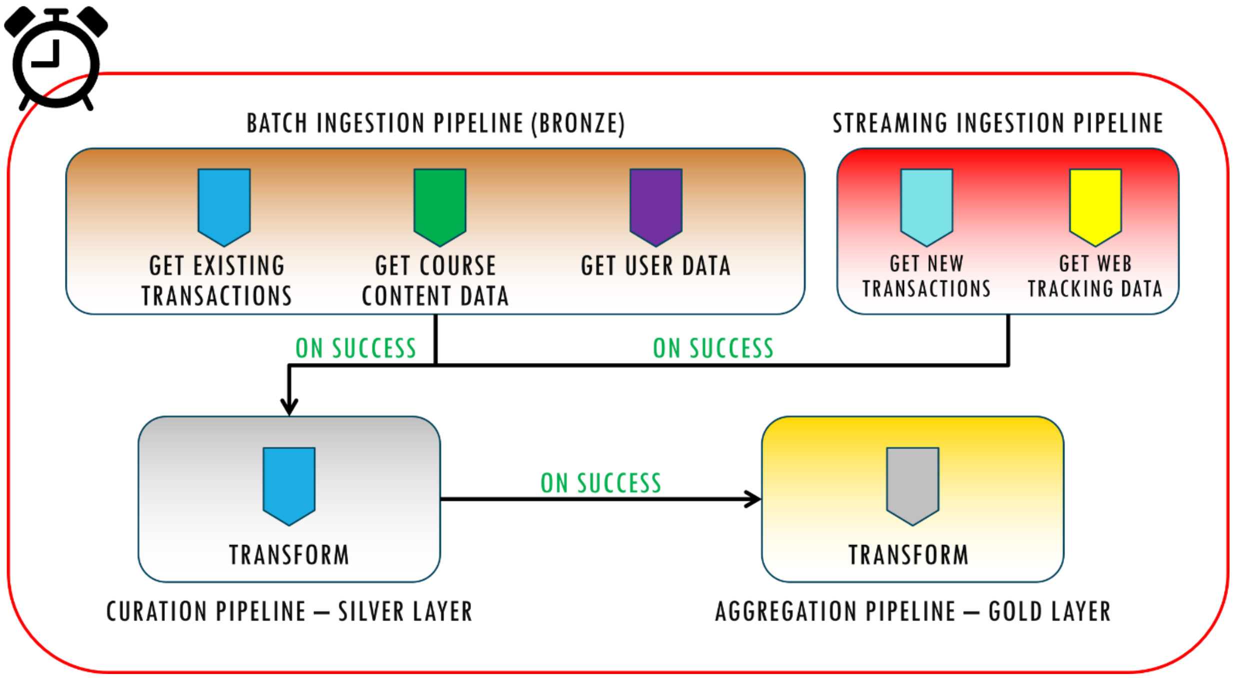

I adopted an **event-based pipeline orchestration** strategy to enable real-time and scheduled data processing without manual intervention.

Key Elements:

- **Master Pipeline**: A controller pipeline designed in **Azure Data Factory** to monitor and trigger sub-pipelines.
- **Sub-Pipelines**: Handle batch ingestion, curation, and aggregation tasks.

Each pipeline is triggered **based on events**, such as:

- Completion or failure of another pipeline
- Arrival of new files in Blob Storage
- Alerts from Azure Monitor

This approach ensures **flexibility and automation**, aligning with the dynamic needs of an online learning platform.

VII.VI PIPELINE FAILURE HANDLING

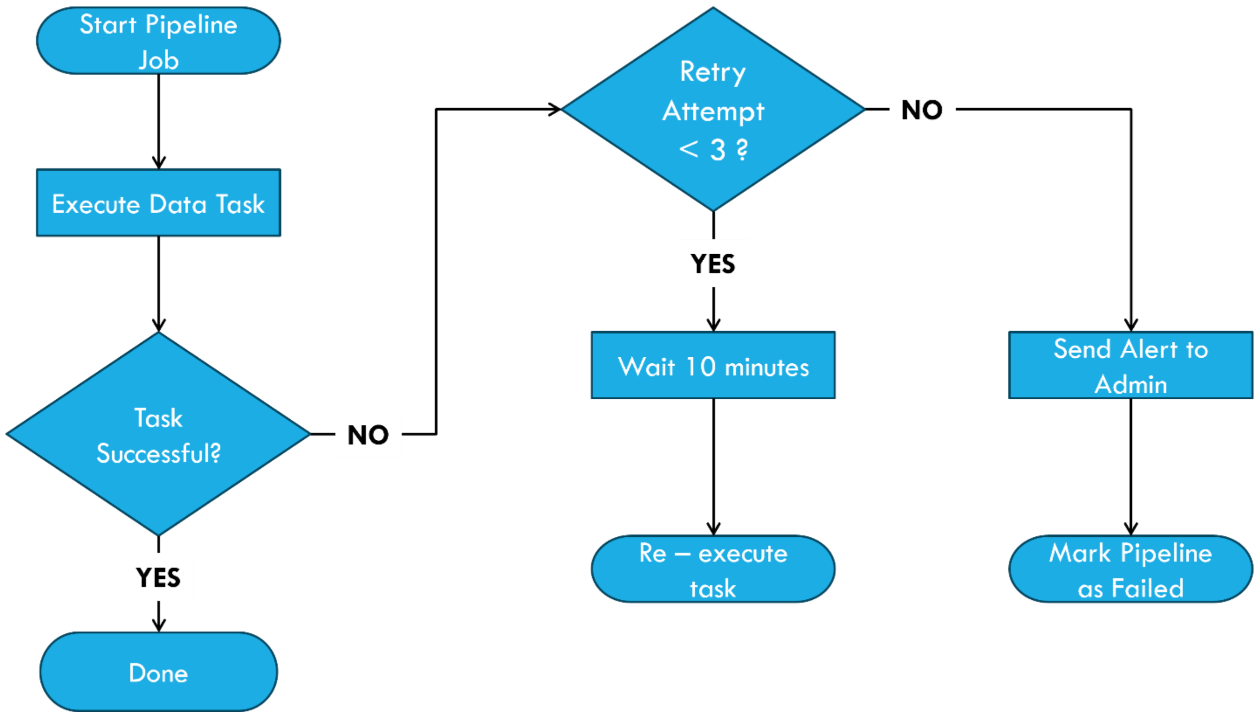

I implemented robust failure handling mechanisms to make pipelines fault-tolerant and self-healing:

Retry Strategy:

- **Retry Attempts:** Each pipeline is configured with **up to 3 retry attempts** upon failure.
- **Retry Interval:** A delay of **10 minutes** between retries to allow for transient issues (e.g., service outages, slow blob upload) to resolve.

If retries still fail, the master pipeline:

- Sends an alert using **Azure Monitor** and **Logic Apps**
- Triggers **incident workflows** (e.g., notifying engineering teams, logging error into tracking systems)

This ensures high pipeline uptime and reduces operational bottlenecks.

VIII. Conclusion

YOULearn.com's data architecture offers a modern, scalable, and resilient approach to handling platform data. By layering raw, curated, and aggregated data, it supports faster decision-making, smarter recommendations, and improved platform performance.

The architecture uses Azure-native tools:

- **Data Factory** for orchestration
- **Event Hub** for real-time streaming
- **Synapse Analytics** for transformation and storage
- **Power BI** for visualization
- **Azure Monitor** for observability
- **Azure ML** for AI-driven personalization

With built-in failure recovery, automation, and machine learning, YOULearn.com is now better equipped to scale, adapt, and innovate in the competitive edtech space.

IX. Future Plans

- Add CI/CD pipelines for deployment automation
- Explore Azure Purview for data cataloging and governance
- Extend ML models to predict student churn
- Enable external data monetization using Azure Data Share
- Introduce role-based access control with Azure Active Directory